



Skinny

a simplified implementation of Chubby

Cody Tseng

Ping-Han Chuang

Wuh-Chwen Hwang

Content

- Design
 - Features
 - RPC call
 - System structure
- Experiment Setup
- Performance Measurement
- Test and Demo

Design – Features

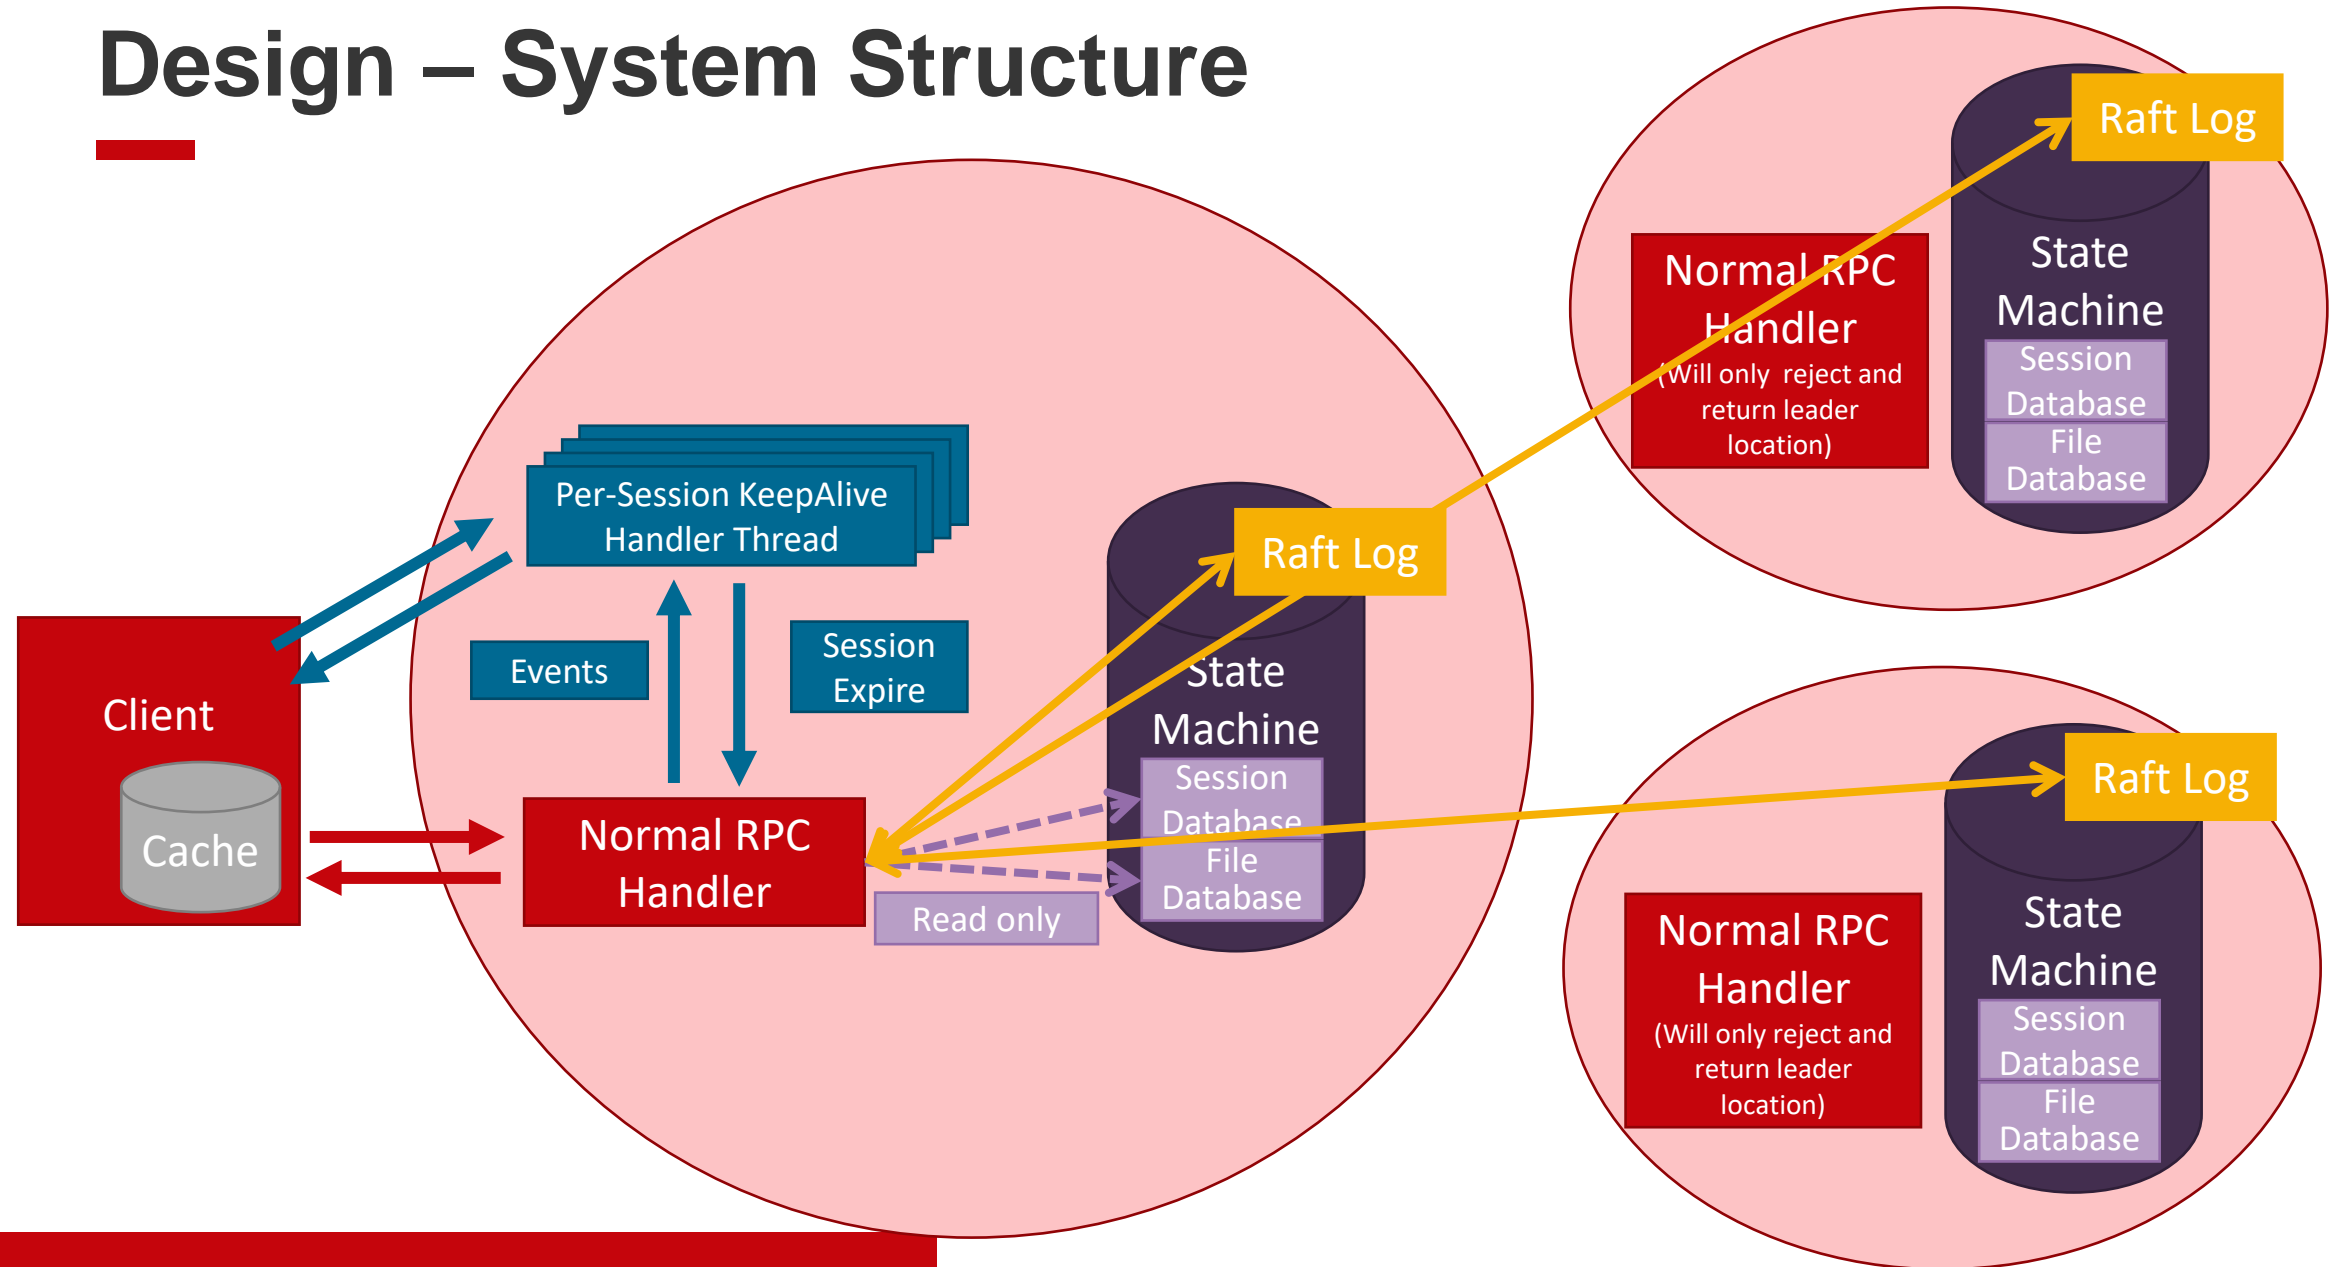
- Aim to be as feature complete as possible
- Supported Features
 - Client session
 - File open/close/delete
 - File read/write
 - Reader/writer lock
 - Single kind of event
 - Directory
 - Client-side Cache
 - Ephemeral files
- Unsupported Features
 - ACL, various kinds of event, snapshot, proxies, lock sequencer and lock upgrading

Design – System Structure

Server

- 3 or 5 servers in a Skinny cell
 - ebay/NuRaft library
 - In-memory states
- Client Library
 - Send keep alive
 - Occasionally receive event and cache invalidation request

Design – System Structure



Design – Files, directories, and handles

- Identify via path name. E.g. /hello/world
- Per file metadata
 - `int instance_num`
 - `bool file_exists`
 - `bool is_directory`
 - `bool is_ephemeral`
 - `bool is_locked_ex`
 - `lock_owners;`
 - `opened_Session;`
- File handle
 - Unlike Chubby, client sees file handle as simply an integer
 - Session DB saves instance number

Design - Locks

- Reader-writer lock
 - Blocking & non-blocking semantics
 - Lock sequencer and lock delay not implemented
-
- `(bool) TryAcquire(fh, ex)`
 - `() Acquire(fh, ex)`
 - `() Release (fh)`

Design - Events

- User can specify event callback when `Open()`
- Piggyback in keepalive response
- Client lib invokes callback when:
 - File: modified / deleted
 - Directory: contained file(s) added / removed
- `(fh) Open (path, callback, is_ephemeral)`
- `(fh) OpenDir (path, callback, is_ephemeral)`


Design – Client API

Comply with Chubby: every API except `Open(Dir)` uses `fh`

- (constructor) [implicitly called `StartSession()`]
- (destructor) [implicitly called `EndSession()`]
- (fh) `Open (path, callback, is_ephemeral)`
- (fh) `OpenDir (path, callback, is_ephemeral)`
- () `Close(fh)`
- (content) `GetContent(fh)`
- () `SetContent(fh, content)`
- () `Delete(fh)`
- (bool) `TryAcquire(fh, ex)`
- () `Acquire(fh, ex)`
- () `Release (fh)`

Design - Caching

- Client will cache data content
- Server invalidate cache when handling write
- Writes will not return until all session acked or expired



Design - Session + KeepAlive

- Client:
 - Must send KeepAlive periodically to maintain an active session.
 - If server does not response in 10s, mark the session as expired
 - Subsequent calls throw error
 - Invalidate cache if KeepAlive returned with a file handle
 - Event: call callback if the user previously register an event associate with that file handle

Design - Sessions + KeepAlive

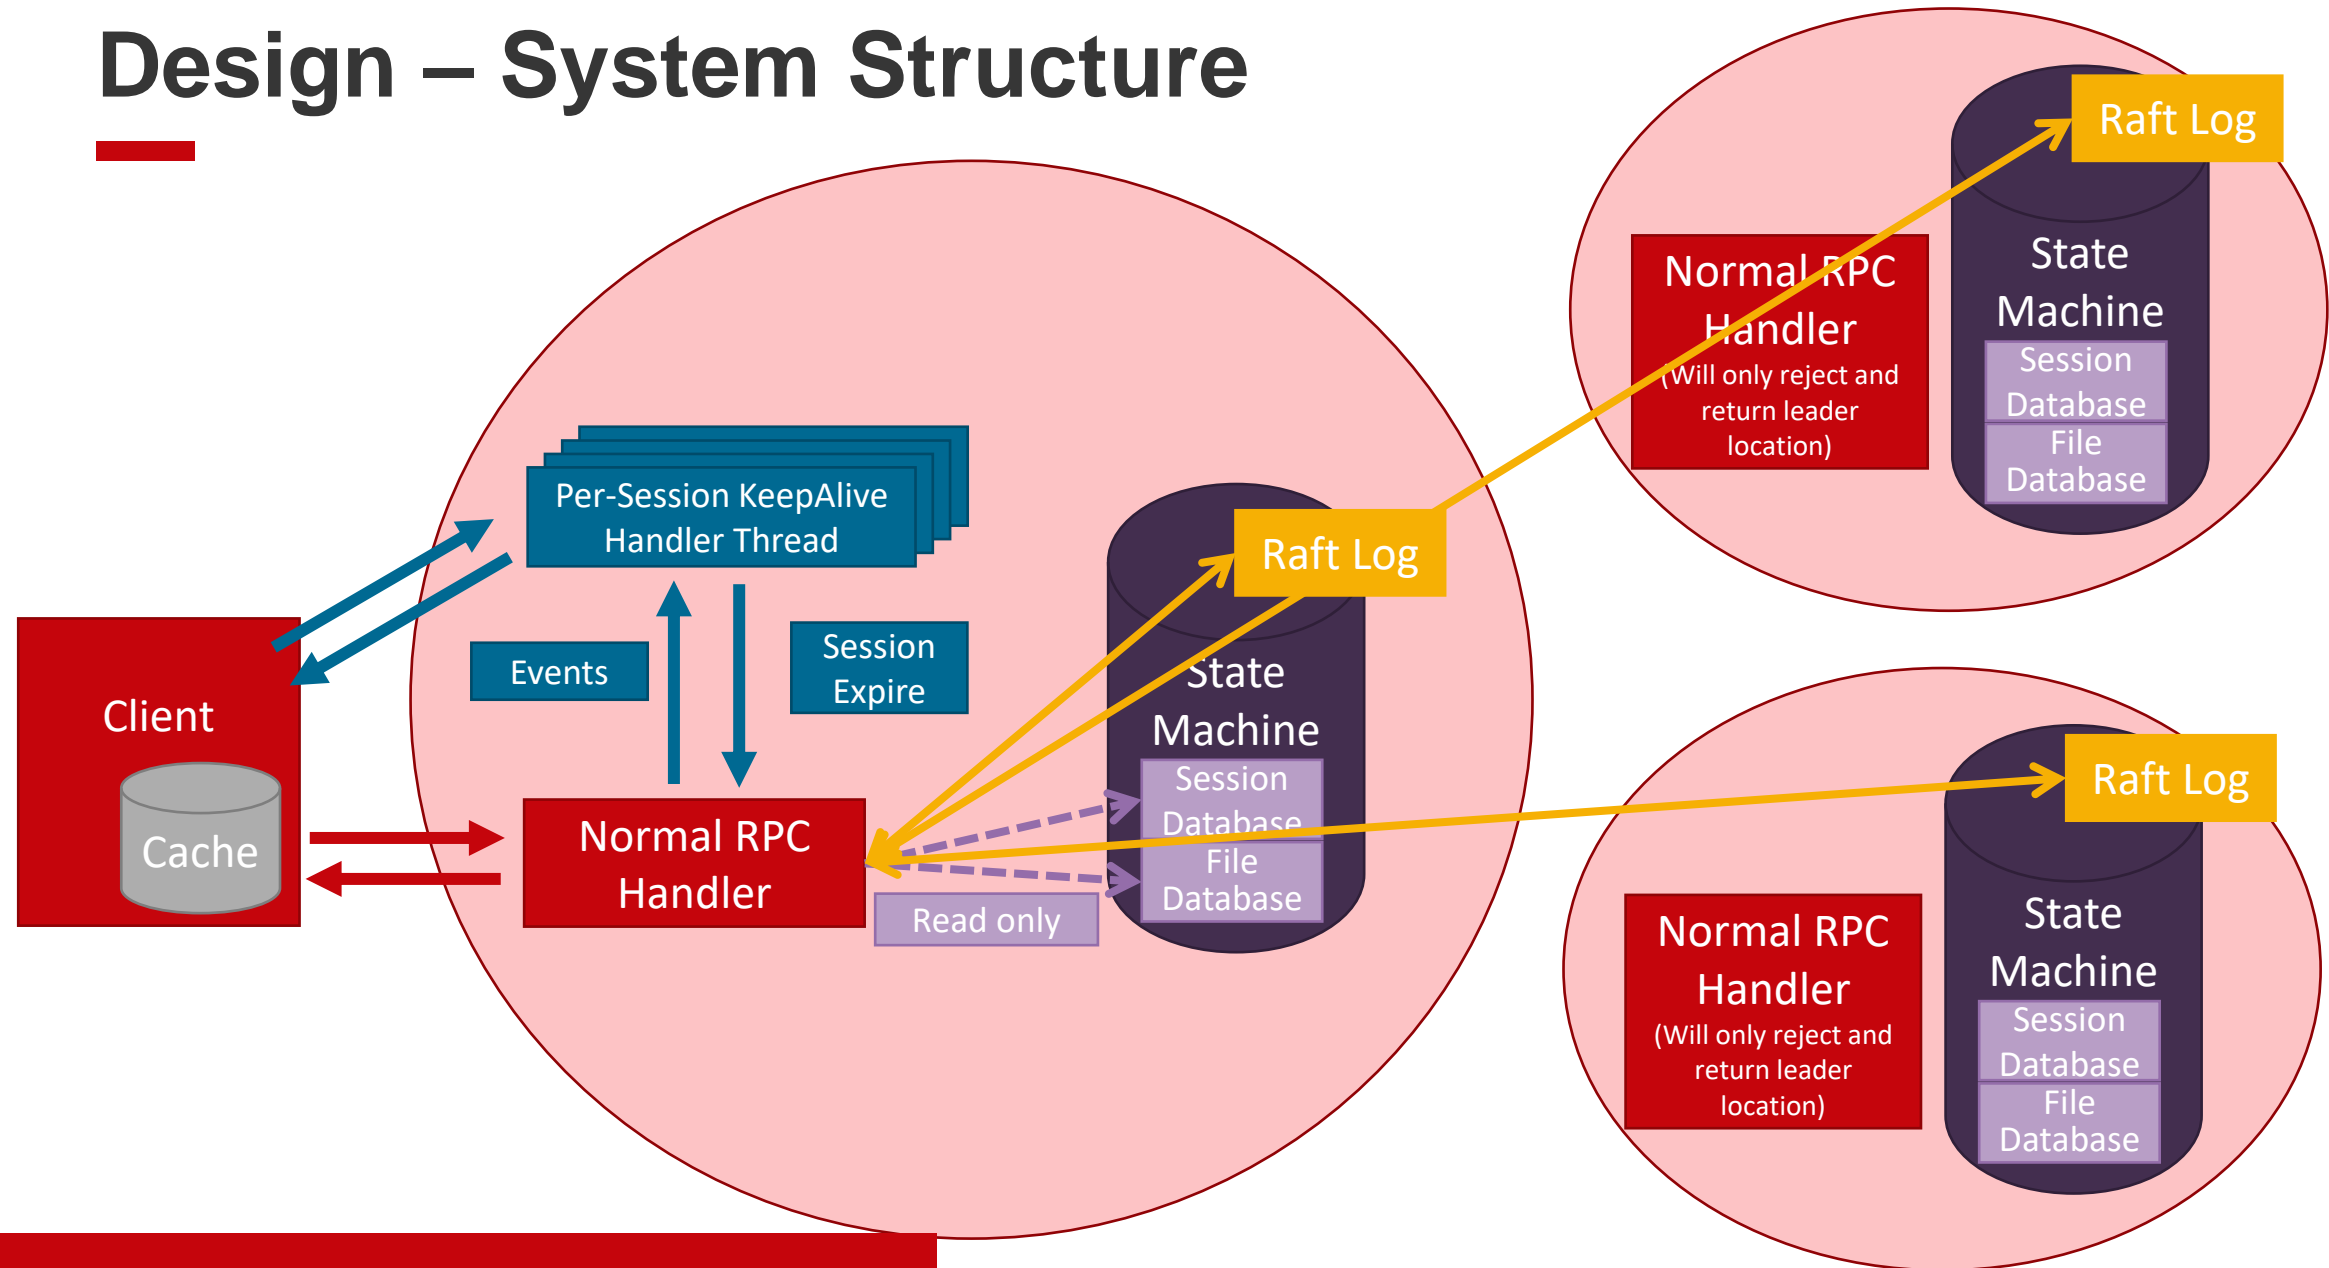
Server: per-session KeepAlive handler thread

- Thread wakes up due to:
 - Receive a new KeepAlive request within timeout => replace current request
 - If there is an event in queue, return immediately w/ message
 - Otherwise, wait 5 seconds, then reset timeout, return call, sleep
 - Receive end_session: cancel and clean up
 - Receive a new event: put in queue
 - Timeout: Kill session
- Timeout: 5s

Leader Failure

- Client – Clear cache, block all requests
- Server – Elect new leader, start sessions' KeepAlive threads (imply timer reset)

Design – System Structure



Correctness Testing

- Bind C++ client library to Python code
- Write test suites in Python (easier to write)
- Able to launch/kill servers in Python
- Test locks/reads/writes with 1000s of clients, events, leader failure, etc.
- Found multiple bugs, mostly race condition

Correctness Testing

- Read/write
 - Single writer, many writes
 - Single writer, one write, many readers
- Lock
 - Single user Acquire, TryAcquire, Release; Acquire after Release
 - Multiple users contending (non-blocking) the same lock
 - Multiple users wait (**blocked**), acquire, and release the same lock
 - **Single writer Acquire** after multiple readers Release
 - **Multiple readers Acquire** after a single writer Release

Correctness Testing

- Event
 - **Event callback** invocation
 - Distributed **barrier** (used in performance measurement)
- Cache
 - Cache invalidation under client failures
- Leader failure

Performance Measurement

Hardware

- Cloudlab c8220 - 40 logical cores, 256 GB RAM

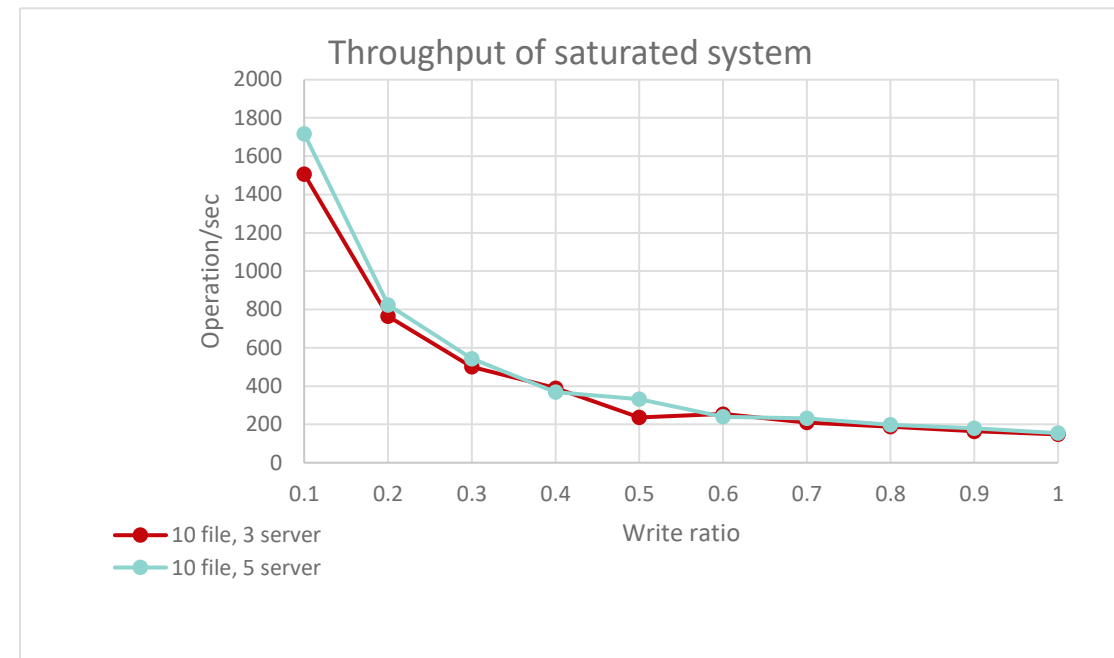

Server: 3/5 nodes cluster

Client (C++)

- 5 nodes, each runs 30 clients
- Open certain amount of files
- Start simultaneously (barrier by Skinny)
- Keep issuing requests to those files for 30 secs
- Each request: randomly [pick a file] & [read or write]
- Parameter:
 - # files
 - Request read/write ratio
- Measurement: throughput (ops/sec), latency


Performance Measurement - Throughput

- Throughput decreases as write ratio increases
 - # reads drop, # writes keep the same
 - Write is the bottleneck
- Number of servers does not affect the result much
 - Servers sit in the same data center



Performance Measurement - Throughput

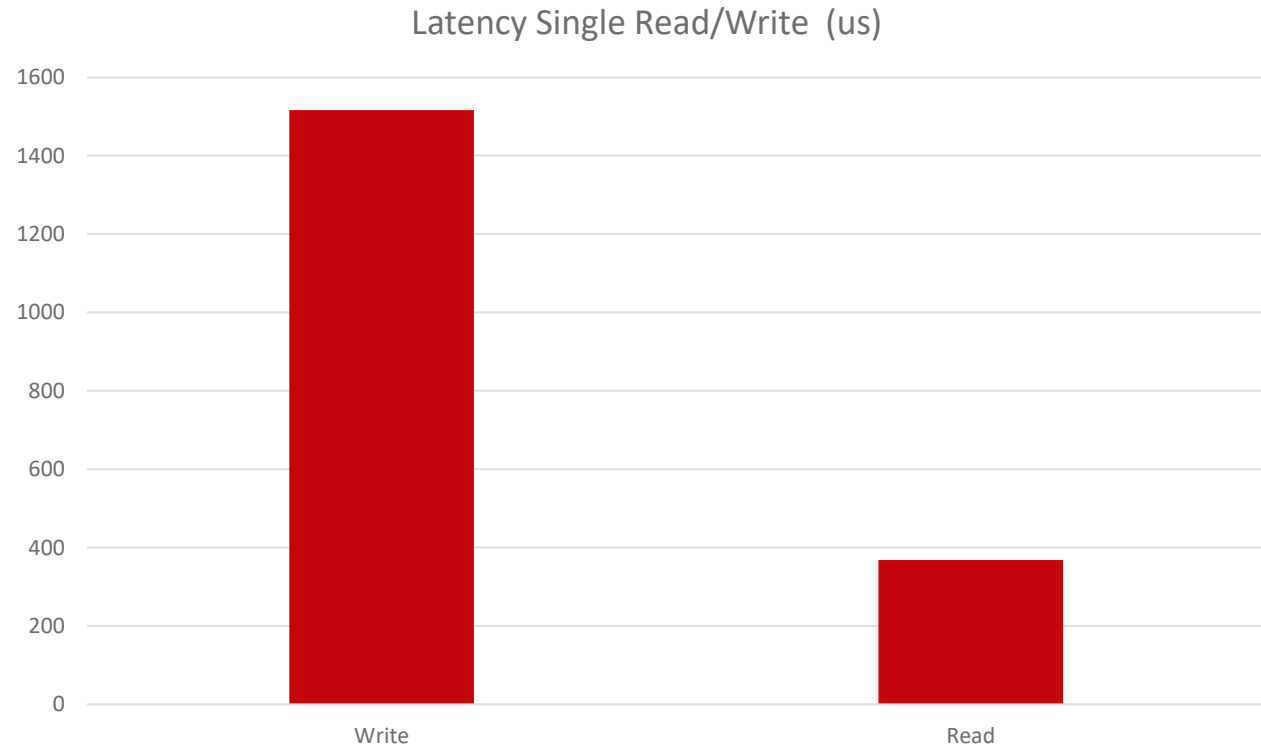
- Changing # opened files
 - Hypothesis: lower # files => higher chance of cache invalidation on the file to be read
 - Result: does not affect much
 - Probably need much more files
 - Opening files takes too much time in Skinny



Performance Measurement – Average Latency

- Single Client
- Single Request

Write	1516.57 us
Read	368.163 us



Demo 1 – Basic functionalities & Leader crash

- Basic functionalities – Open / Read / Write / Lock
- Show that previous content can still be read even if the leader crash
- Show that lock status is still maintained after a leader crash

Demo 2 – Membership changes

```
SkinnyClient client = SkinnyClient();
int fh = client.OpenDir("/service", [&client](int fh) {
    /* Print directory content */
});

bool is_ephemeral = true;
client.Open("/service/" + std::string(hostname), std::nullopt, is_ephemeral);
```

Demo 3 – Primary Election

```
SkinnyClient a = SkinnyClient();
int fh = a.Open("/primary", [&a](int fh) {
    std::cout << "The primary server is " << a.GetContent(fh) << std::endl;
});

std::string primary = a.GetContent(fh);

if (primary.empty()) { // If the primary server had already be selected
    bool acq_success = a.TryAcquire(fh, true); // try to get an exclusive lock
    if (acq_success) {
        a.SetContent(fh, hostname);
    }
} else {
    std::cout << "The primary server is " << primary << std::endl;
}
```




Lessons Learned

- Should've clearly defined where to put logic
 - State Machine vs GRPC Service
- State Machine operations should never block
- Should've used some library that provide synchronized container
 - `Synchronized<vector<int>> vec;`
 - `vec.wlock()->push_back(3)`